

RUBAB ZAHRA SARFRAZ

rubabzsarfraz@gmail.com · rubabzsarfraz.com · linkedin.com/in/rubabzsarfraz

SUMMARY

I'm a data science professional with over 6 years of experience in building data science teams and engineering data-driven solutions that serve over half a million users. My industry experience of solving real-world data challenges has motivated me to pursue a PhD focused on databases in which I can study data systems in depth and help build efficient and responsible human-database interactions.

EDUCATION

Lahore University of Management Sciences (LUMS)

Aug. 2016 – Jun. 2018

M.S. in Computer Science

Lahore, Pakistan

- Thesis: Measuring the Impact of Fake News in Developing Regions
- Advisor: [Dr. Ihsan Ayyub Qazi](#)

University of Engineering & Technology (UET)

Oct. 2012 – Jun. 2016

B.Sc. in Computer Engineering (with honors)

Lahore, Pakistan

- Thesis: Monitoring Traffic on Virtual Routers of OpenStack
- Advisor: [Dr. Irfan Ullah Chaudhary](#)
- GPA: 3.74/4.00

PUBLICATIONS

Rubab Zahra Sarfraz, Samar Haider. “Vizard: Improving Visual Data Literacy with Large Language Models”. In *7th International Workshop on Big Data Visual Exploration and Analytics (BigVis) at VLDB 2024*. [[Slides](#)][[Code](#)][[Video](#)] [[Demo](#)]

Rubab Zahra Sarfraz. “Towards Semi-Supervised Data Quality Detection in Graphs”. In *13th International Workshop on Quality in Databases (QDB) at VLDB 2024*. [[Slides](#)][[Code](#)][[Video](#)]

Nida Munawar, **Rubab Zahra Sarfraz**, Maria Costello, David Robinson, Colm Bergin, Elaine Greene. “Risk Factors and Outcomes of Delirium in Hospitalized Older Adults with COVID-19: A Systematic Review and Meta-Analysis”. In *Aging and Health Research (2023)*.

WORK EXPERIENCE

BridgeLinx

Oct. 2022 – Present

Data Lead

Lahore, Pakistan

- Established the company's data infrastructure from the ground up, utilizing Snowflake for data lakes, Prefect for quality-driven data pipelines, and Tableau/Metabase for front-end, adopted by **40%** of the workforce.
- Built a temporal geospatial data pipeline, boosting customer acquisition by **25%** across **150+ transportation lanes** through optimized operational reach and resource allocation.
- Incorporated data quality checklist in product sprints. Created and deployed a data quality bot, improving data completeness from **45% to 98%**, resulting in 30% faster decision-making across **9 teams** and **200+ data points**.
- Engineered a real-time bidding system employing advanced analytics to optimize Return on Capital (ROC) through predictive client settlement behavior analysis upon order booking.

Finja

Jan. 2020 – Oct. 2022

Data Lead

Lahore, Pakistan

- Researched the low-resourced, undocumented SME market, engineered a machine learning pipeline identifying 40+ predictive variables, reducing loan approval time by **35%**, while keeping the **NPLs below 0.5%** with **\$3M issuances/month**.
- Advocated for and expanded the ML model's business evaluation metric from overall to category-wise NPLs to protect vulnerable MSMEs from market-driven penalties, using a real-time feedback loop for classifying bad loans into categories.
- Led a multidisciplinary team to launch Pakistan's first data-driven **SECP-approved** P2P lending, investment and payment **engine**, resulting in **\$1M+** lending issuance to **3000+ retail stores**.

- Led a 10-member data science team, achieving a **40%** reduction in data processing time and transitioning 8 departments to fully adopt data-driven decision-making, increasing operational efficiency by **25%**.
- Instrumental in securing a **\$9M Series A** funding by managing the data rooms, drafting and presenting the impact of data-driven tech and insights to investors.

Oct. 2018 – Dec. 2019

Data Engineer

Lahore, Pakistan

- Architected a scalable data architecture for financial transactions (lending, payments, investments), optimizing for performance, security, and compliance across a multi-service platform serving **300K customers**.
- Built and managed a high-availability data lake and hybrid data pipelines, integrating real-time streaming and batch processing to prioritize workloads efficiently across **50+ dashboards**, **8 departments** and **3 products**.
- Developed a high-accuracy sanctions screening system in-house for KYC, enabling real-time fraud detection on incoming wallets and saving **\$150K**. Successfully identified **two high-risk matches**, mitigating compliance risks and preventing financial losses.
- Enhanced search engine relevance for a B2C marketplace (from **40%** to **80%** relevance ratio) by improving vendor product data quality and using machine learning models in an adaptive pipeline with human-in-the-loop validation.

CERN

Jul. 2018 – Sept. 2018

Intern, Software Engineer

Geneva, Switzerland

- Spearheaded the deployment of **Ceph clusters with Rook on Kubernetes**, leveraging both OpenStack VMs and Ironic hosts.
- Defined and implemented robust evaluation metrics, demonstrating the superior latency and user-friendliness of the new cluster deployment approach, reducing setup time by **50%**.
- Enhanced the Orchestrator CLI with **RGW support**, enabling rapid S3 service provisioning in just **1-2 minutes**.
- Published a **blog post** with the findings in collaboration with CERN and Ceph (RedHat).

RedHat (via Outreachy)

Dec. 2017 – Mar. 2018

Intern, Software Engineer

Remote

- Improved distributed cluster management by adding performance dashboards. Pull requests: [1] [2] [3]
- Converted the decentralized architecture of Ceph Manager to a **centralized one** for improved coherence with the codebase.

Meta

Feb. 2017 – Mar. 2017

Mentee, Software Engineer

Remote

- Extended the Osquery open source project by implementing a **virtual table in C++ for listing Python packages** installed on a server.
- This functionality was actively used for securing Meta's data center servers against **vulnerabilities introduced by PyPy in 2017**.

TEACHING EXPERIENCE

Instructor , Introduction to Data Science Lahore School of Economics	Summer 2023
Teaching Assistant , Advanced Operating Systems Lahore University of Management Sciences Instructor: Dr. Muhammad Hamad Alizai	Spring 2017
Teaching Assistant , Programming Fundamentals University of Engineering & Technology Instructor: Dr. Irfan Ullah Chaudhary	Spring 2016

AWARDS & HONORS

Invited Participant, International Visitor Leadership Program (IVLP), U.S. Department of State	2024
Finalist, USAID FDI Grant of \$100K (on behalf of BridgeLinx)	2024
Finalist, U.K. Climate Finance Accelerator (on behalf of BridgeLinx)	2023
Runner Up, CERN Openlab Lightning Talks	2018
Diversity Scholar, KubeCon by the Linux Foundation	2018
Dean's Honor List, University of Engineering & Technology	2012 – 2016

COMMUNITY INVOLVEMENT

Reviewer, ACM Conference on Human Factors in Computing Systems (CHI)	2025
Reviewer, IEEE VIS Workshop on Visualization for AI Explainability (VISxAI)	2024
Board Member, Computer Science Alumni Network, LUMS	2024
Advisor, Securities and Exchange Commission of Pakistan (for drafting regulations for P2P lending)	2022
Open Source Contributor, Ceph Blog: Evaluating Ceph Deployments with Rook	2018
Open Source Contributor, Docker Docs Hackathon. Pull requests: [1] [2] [3] [4]	2017

TALKS

“Elevating Trust in Your Data with Python” , PyCon Pakistan	2024
“Reshaping the Financial Sector with Artificial Intelligence”, Information Technology University	2022
“Debugging the Startup Space: Tech Careers”, LUMS Women in Computing	2022
“Best Practices in Operationalizing Machine Learning” , AWS Startup CTO Forum	2022
“Breaking the Bias”, Systems Limited	2021
“How Technology can be Your Best Friend”, Meta Developers Circle	2020
“Evaluating Ceph Deployment with Rook” , CERN	2018

TECHNICAL SKILLS

Expertise: Data Engineering, Data Architecture, Data Modeling, Data Product Management

Languages & Scripts: Python, C++, SQL, HTML, \LaTeX

Machine Learning: Scikit-learn, PyTorch, Transformers

Data Lake & Quality: BigQuery, Snowflake, S3, Great Expectations, Pandera

Data Visualization: Tableau, Looker Studio, Streamlit

Data Infrastructure: GCP (Pub/Sub, Cloud Functions, Storage, Dataflow, SQL), AWS (SageMaker, Lambda), Prefect, FastAPI, Flask, Docker